

```
In [ ]: import time

import numpy
import matplotlib.pyplot as plt
import numpy as np
import scipy.stats as sp_spec
import scipy.stats as sp_stats
```

## Assignment 1A. Problem 1.4.19 SVI.

### Generate data

The cell below generates data for the LDA model. Note, for simplicity, we are using  $N, d = N$  for all  $d$ .

```
In [ ]: def generate_data(D, N, K, W, eta, alpha):
    # sample K topics
    beta = sp_stats.dirichlet(eta).rvs(size=K) # size K x W
    theta = np.zeros((D, K)) # size D x K

    w = np.zeros((D, N, W))
    z = np.zeros((D, N), dtype=int)
    for d in range(D):
        # sample document topic distribution
        theta_d = sp_stats.dirichlet(alpha).rvs(size=W)
        theta[d, :] = theta_d
        for n in range(N):
            # sample word to topic assignment
            z[d, n] = np.argmax(np.dot(theta_d, beta))
            # sample word
            w[d, n] = sp_stats.multinomial(n=1, p=beta[z[d, n], :]).rvs(1)

    z[d, n] = z[d, n].nd
    w[d, n] = w[d, n].nd

    return w, z, theta, beta

D_sim = 500
N_sim = 50
K_sim = 5

eta_sim = np.ones(W_sim)
eta_sim[3] = 0.0001 # Expect word 3 to not appear in data
eta_sim[1] = 3. # Expect word 1 to be most common in data
alpha_sim = np.ones(K_sim) * 1.0
w0, z0, theta0, beta0 = generate_data(D_sim, N_sim, K_sim, W_sim, eta_sim, alpha_sim)
w_cat = w0.argmax(axis=-1) # remove one not encoding
unique_z, counts_z = numpy.unique(z0(0, :), return_counts=True)
unique_w, counts_w = numpy.unique(w_cat(0, :), return_counts=True)

# Sanity checks for data generation
print(f"Average z of of document should be close to theta of document. \n Theta of doc 0: {theta0[0]} \n Mean z of doc 0: {counts_z[N_sim]}"
print(f"beta of topic 0: {beta0[0]}"
print(f"beta of topic 1: {beta0[1]}"
print(f"Word to topic assignment, z, of document 0: {z0[0, 0:10]}"
print(f"Observed words, w, of document 0: {w_cat(0, 0:10]}"
print(f"Unique words and count of document 0: {[(u, c)] for u, c in zip(unique_w, counts_w)]}"

Average z of each document should be close to theta of document.
Theta of doc 0: [0.54132269 0.45867731]
Mean z of doc 0: [0.5 0.5]
Beta of topic 0: [0.11439051 0.69431184 0.08117848 0. 0.11029637]
Beta of topic 1: [0.18644559 0.37789851 0.47896521 0. 0.03659869]
Word to topic assignment, z, of document 0: [1 2 0 1 0 1 0 1 0 1 0 1]
Observed words, w, of document 0: [1 2 4 1 1 1 1 1 1 2]
Unique words and count of document 0: ['0': 9, '1': 27, '2': 11, '4': 3]
```

```
In [ ]: import torch
import torch.distributions as t_dist

def generate_data_torch(D, N, K, W, eta, alpha):
    """
    Torch implementation for generating data using the LDA model. Needed for sampling larger datasets.
    """
    # sample K topics
    beta_dist = t_dist.Dirichlet(torch.from_numpy(eta))
    beta = beta_dist.sample([K]) # size K x W

    # sample document topic distribution
    theta_dist = t_dist.Dirichlet(torch.from_numpy(alpha))
    theta = theta_dist.sample([D])

    # sample word to topic assignment
    z_dist = t_dist.Categorical(probs=theta)
    z = z_dist.sample([N]) # size D x K

    # sample word from selected topics
    beta_select = torch.einsum("kw, dnk -> dnw", beta, z)
    w_dist = t_dist.Categorical(probs=beta_select)
    w = w_dist.sample([N])

    w = w.reshape(D, N, W)

    return w.numpy(), z.numpy(), theta.numpy(), beta.numpy()
```

### Helper functions

```
In [ ]: def log_multivariate_beta_function(a, axis=None):
    return np.sum(sp_spec.gammaln(a)) - sp_spec.gammaln(np.sum(a, axis=axis))
```

### CAVI Implementation, ELBO and initialization

```
In [ ]: def initialize_q(w, D, N, K, W):
    """
    Random initialization.
    """
    phi_init = np.random.random(size=(D, N, K))
    phi_init = phi_init / np.sum(phi_init, axis=-1, keepdims=True)
    gamma_init = np.random.random((1, 10, size=(D, K)))
    lambda_init = np.random.random((1, 10, size=(K, W)))
    return phi_init, gamma_init, lambda_init

def update_q_Z(w, gamma, lambda):
    D, N, W = w.shape
    K, W = lambda.shape
    W = eta.shape
    E_log_theta = sp_spec.digamma(gamma) - sp_spec.digamma(np.sum(gamma, axis=-1, keepdims=True)) # size D x K
    E_log_beta = sp_spec.digamma(lambda) - sp_spec.digamma(np.sum(lambda, axis=-1, keepdims=True)) # size K x W
    log_rho = np.zeros((D, N, K))
    w_label = w.argmax(axis=-1)
    for d in range(D):
        for n in range(N):
            E_log_beta_wdn = E_log_beta[:, int(w_label[d, n])]
            E_log_theta_d = E_log_theta[d, :]
            log_rho_dn = E_log_theta_d + E_log_beta_wdn
            log_rho[d, n, :] = log_rho_dn

    phi = np.exp(log_rho - sp_spec.logsumexp(log_rho, axis=-1, keepdims=True))
    return phi

def update_q_theta(phi, alpha):
    E_Z = phi
    D, N, K = phi.shape
    gamma = np.zeros((D, K))
    for d in range(D):
        E_Z_d = E_Z[d]
        gamma[d, :] = alpha + np.sum(E_Z_d, axis=0) # sum over N
    return gamma

def update_q_beta(w, phi, eta):
    E_Z = phi
    D, N, W = w.shape
    K = phi.shape[-1]
    lambda = np.zeros((K, W))
    for k in range(K):
        lambda[k, :] = eta
        for d in range(D):
            for n in range(N):
                lambda[k, :] += E_Z[d, n, k] * w[d, n] # Sum over d and n
    return lambda

def calculate_elbo(w, phi, gamma, lambda, eta, alpha):
    D, N, K = phi.shape
    W = eta.shape
    E_log_theta = sp_spec.digamma(gamma) - sp_spec.digamma(np.sum(gamma, axis=-1, keepdims=True)) # size D x K
    E_log_beta = sp_spec.digamma(lambda) - sp_spec.digamma(np.sum(lambda, axis=-1, keepdims=True)) # size K x W
    E_Z = phi # size D, N, K
    log_Beta_alpha = log_multivariate_beta_function(alpha)
    log_Beta_eta = log_multivariate_beta_function(eta)
    log_Beta_gamma = np.array([log_multivariate_beta_function(gamma[d, :]) for d in range(D)])
    dg_gamma = sp_spec.digamma(gamma)
    log_Beta_lambda = np.array([log_multivariate_beta_function(lambda[k, :]) for k in range(K)])
    dg_lambda = sp_spec.digamma(lambda)

    neg_CE_likelihood = np.einsum("dnk, kw, dnw", E_Z, log_Beta_eta, w)
    neg_CE_Z = np.einsum("dnk, dk -> ", E_Z, E_log_theta)
    neg_CE_theta = -D * log_Beta_alpha - np.einsum("k, dk -> ", alpha - 1, E_log_theta)
    neg_CE_beta = -K * log_Beta_eta - np.einsum("w, kw -> ", eta - 1, E_log_beta)
    H_Z = -np.einsum("dnk, dnk -> ", E_Z, np.log(E_Z))
    dg_gamma = sp_spec.digamma(gamma)
    H_theta = np.sum(log_Beta_gamma + (gamma - 0 - K) * dg_gamma - np.einsum("dk, dk -> d", gamma - 1, dg_gamma))
    lambda_0 = np.sum(lambda, axis=-1)
    dg_lambda0 = sp_spec.digamma(lambda_0)
    H_beta = np.sum(log_Beta_lambda + (lambda_0 - W) * dg_lambda0 - np.einsum("kw, kw -> k", lambda - 1, dg_lambda))
    return neg_CE_likelihood + neg_CE_Z + neg_CE_theta + neg_CE_beta + H_Z + H_theta + H_beta

def CAVI_algorithm(w, K, N_iter, eta, alpha):
    D, N, W = w.shape
    phi, gamma, lambda = initialize_q(w, D, N, K, W)

    # Store output per iteration
    elbo = np.zeros(N_iter)
    phi_out = np.zeros((N_iter, D, N, K))
    gamma_out = np.zeros((N_iter, D, K))
    lambda_out = np.zeros((N_iter, K, W))

    for i in range(0, N_iter):
        ##### CAVI updates #####

        # q(Z) update
        phi = update_q_Z(w, gamma, lambda)

        # q(theta) update
        gamma = update_q_theta(phi, alpha)

        # q(beta) update
        lambda = update_q_beta(w, phi, eta)

        # ELBO
        elbo[i] = calculate_elbo(w, phi, gamma, lambda, eta, alpha)

    # Outputs
    phi_out[1:] = phi
    gamma_out[1:] = gamma
    lambda_out[1:] = lambda

    return phi_out, gamma_out, lambda_out, elbo

N_iter0 = 100
K0 = K_sim
W0 = W_sim
eta_prior0 = np.ones(W0)
alpha_prior0 = np.ones(K0)
phi_out0, gamma_out0, lambda_out0, elbo0 = CAVI_algorithm(w0, K0, N_iter0, eta_prior0, alpha_prior0)
final_phi0 = phi_out0[-1]
final_gamma0 = gamma_out0[-1]
final_lambda0 = lambda_out0[-1]
```

```
In [ ]: precision = 3
print(f"----- Recall label switching - compare E[theta] and true theta and check for label switching -----")
print(f"Final E[theta] of doc 0 CAVI: (np.round(final_gamma0[0], axis=0, keepdims=True), precision)")
print(f"True theta of doc 0: (np.round(theta0[0], precision)")

print(f"----- Recall label switching - e.g. E[beta_0] could be fit to true theta.1. -----")
print(f"Final E[beta] k=1: (np.round(final_lambda0[1, :]) / np.sum(final_lambda0[1, :]), axis=-1, keepdims=True), precision)")
print(f"Final E[beta] k=1: (np.round(beta0[1, :]) / np.sum(beta0[1, :]), axis=-1, keepdims=True), precision)")
print(f"True beta k=1: (np.round(beta0[1, :]), precision)")

----- Recall label switching - compare E[theta] and true theta and check for label switching -----
Final E[theta] of doc 0 CAVI: [0.394 0.606]
True theta of doc 0: [0.541 0.459]
----- Recall label switching - e.g. E[beta_0] could be fit to true theta.1. -----
Final E[beta] k=0: [0.117 0.249 0.623 0. 0.011]
Final E[beta] k=1: [0.187 0.378 0.484 0. 0.120]
True beta k=0: [0.114 0.694 0.881 0. 0.11]
True beta k=1: [0.186 0.378 0.479 0. 0.037]
```

### SVI Implementation

Using the CAVI updates as a template, finish the code below.

```
In [ ]: def update_q_Z_svi(batch, w, gamma, lambda):
    """
    SVI update for q(Z).
    """
    D_batch, N_batch, _ = batch.shape
    K = gamma.shape[-1]

    E_log_theta = sp_spec.digamma(gamma) - sp_spec.digamma(np.sum(gamma, axis=-1, keepdims=True))
    E_log_beta = sp_spec.digamma(lambda) - sp_spec.digamma(np.sum(lambda, axis=-1, keepdims=True))

    log_rho = np.zeros((D_batch, N_batch, K))
    w_label = batch.argmax(axis=-1)

    for d in range(D_batch):
        for n in range(N_batch):
            E_log_beta_wdn = E_log_beta[:, int(w_label[d, n])]
            E_log_theta_d = E_log_theta[d, :]
            log_rho_dn = E_log_theta_d + E_log_beta_wdn
            log_rho[d, n, :] = log_rho_dn

    phi_batch = np.exp(log_rho - sp_spec.logsumexp(log_rho, axis=-1, keepdims=True))
    return phi_batch

def update_q_theta_svi(batch, phi, alpha):
    """
    SVI update for q(theta).
    """
    D_batch, _, K = phi.shape

    E_Z_batch = phi
    gamma_batch = np.zeros((D_batch, K))

    for d in range(D_batch):
        E_Z_d = E_Z_batch[d]
        gamma_batch[d, :] = alpha + np.sum(E_Z_d, axis=0) # Sum over N
    return gamma_batch

def update_q_beta_svi(batch, w, phi, eta):
    """
    SVI update for q(beta).
    """
    D_batch, N_batch, W = batch.shape
    K = phi.shape[-1]

    lambda_batch = np.zeros((K, W))

    for k in range(K):
        lambda_batch[k, :] = eta
        for d in range(D_batch):
            for n in range(N_batch):
                lambda_batch[k, :] += phi[d, n, k] * batch[d, n] # Sum over d and n
    return lambda_batch

def SVI_algorithm(w, K, S, N_iter, eta, alpha):
    """
    Stochastic Variational Inference (SVI) algorithm for LDA.
    """
    D, N, W = w.shape
    phi, gamma, lambda = initialize_q(w, D, N, K, W)

    # Store output per iteration
    elbo = np.zeros(N_iter)
    phi_out = np.zeros((N_iter, D, N, K))
    gamma_out = np.zeros((N_iter, D, K))
    lambda_out = np.zeros((N_iter, K, W))

    for i in range(0, N_iter):
        # Sample batch (subsample documents)
        batch_indices = np.random.choice(D, size=S, replace=False)
        batch_w = w[batch_indices]

        # SVI updates
        phi_batch = update_q_Z_svi(batch_w, w, gamma, lambda)
        gamma_batch = update_q_theta_svi(batch_w, phi_batch, alpha)
        lambda_batch = update_q_beta_svi(batch_w, w, phi_batch, eta)

        # Update global variables
        phi[batch_indices, :] = phi_batch
        gamma[batch_indices, :] = gamma_batch
        lambda[batch_indices, :] = lambda_batch

        # ELBO
        elbo[i] = calculate_elbo(w, phi, gamma, lambda, eta, alpha)

    # Store outputs
    phi_out[1:] = phi
    gamma_out[1:] = gamma
    lambda_out[1:] = lambda

    return phi_out, gamma_out, lambda_out, elbo
```

### CASE 1

#### Tiny dataset

```
In [ ]: np.random.seed(0)

# Data simulation parameters
D1 = 50
K1 = 50
K1 = 2
W1 = 5
eta_sim1 = np.ones(W1)
alpha_sim1 = np.ones(K1)

w1, z1, theta1, beta1 = generate_data(D1, N1, K1, W1, eta_sim1, alpha_sim1)

# Inference parameters
N_iter_cavi1 = 100
N_iter_svi1 = 100
eta_prior1 = np.ones(W1) * 1.
alpha_prior1 = np.ones(K1) * 1.
S1 = 5 # batch size

start_cavi1 = time.time()
phi_out1_cavi, gamma_out1_cavi, lambda_out1_cavi, elbo1_cavi = CAVI_algorithm(w1, K1, N_iter_cavi1, eta_prior1, alpha_prior1)
end_cavi1 = time.time()

start_svi1 = time.time()
phi_out1_svi, gamma_out1_svi, lambda_out1_svi, elbo1_svi = SVI_algorithm(w1, K1, S1, N_iter_svi1, eta_prior1, alpha_prior1)
end_svi1 = time.time()

final_phi1_cavi = phi_out1_cavi[-1]
final_gamma1_cavi = gamma_out1_cavi[-1]
final_lambda1_cavi = lambda_out1_cavi[-1]
final_phi1_svi = phi_out1_svi[-1]
final_gamma1_svi = gamma_out1_svi[-1]
final_lambda1_svi = lambda_out1_svi[-1]
```

#### Evaluation


Do not expect perfect results in terms of expectations being identical to the "true" theta and beta. Do not expect the ELBO plot of your SVI alg to be the same as the CAVI alg. However, it should increase and be in the same ball park as that of the CAVI alg.

```
In [ ]: np.set_printoptions(formatter={"float": lambda x: "{0:0.3f}".format(x)})
print(f"----- Recall label switching - compare E[theta] and true theta and check for label switching -----")
print(f"Final E[theta] of doc 0 SVI: (final_gamma1_svi[0] / np.sum(final_gamma1_svi[0], axis=0, keepdims=True)")
print(f"True theta of doc 0: (theta0[0], precision)")

print(f"----- Recall label switching - e.g. E[beta_0] could be fit to true theta.1. -----")
print(f"Final E[beta] k=0: (final_lambda1_svi[0, :]) / np.sum(final_lambda1_svi[0, :]), axis=-1, keepdims=True)")
print(f"Final E[beta] k=1: (final_lambda1_svi[1, :]) / np.sum(final_lambda1_svi[1, :]), axis=-1, keepdims=True)")
print(f"True beta k=0: (beta0[0, :]) / np.sum(beta0[0, :]), axis=-1, keepdims=True)")
print(f"True beta k=1: (beta0[1, :]) / np.sum(beta0[1, :]), axis=-1, keepdims=True)")

----- Recall label switching - compare E[theta] and true theta and check for label switching -----
E[theta] of doc 0 SVI: [0.686 0.314]
E[theta] of doc 0 CAVI: [0.475 0.525]
True theta of doc 0: [0.541 0.459]
----- Recall label switching - e.g. E[beta_0] could be fit to true theta.1. -----
E[beta] k=0: [0.142 0.858 0.083 0.145 0.027 0.058 0.018 0.183 0.161 0.093]
E[beta] k=1: [0.152 0.858 0.087 0.144 0.023 0.067 0.021 0.181 0.163 0.096]
E[beta] k=2: [0.164 0.866 0.083 0.144 0.030 0.069 0.021 0.184 0.177 0.071]
True beta k=0: [0.067 0.185 0.077 0.060 0.046 0.087 0.040 0.186 0.177 0.040]
True beta k=1: [0.138 0.867 0.074 0.138 0.087 0.068 0.002 0.158 0.134 0.181]
True beta k=2: [0.295 0.123 0.047 0.116 0.010 0.078 0.012 0.222 0.057 0.041]
Time CAVI: 0.015250509943
Time SVI: 44.14759862121582
```

```
In [ ]: plt.plot(list(range(1, N_iter_cavi1 + 1)), elbo1_cavi(np.arange(0, N_iter_svi1, int(N_iter_svi1 / N_iter_cavi1))))
plt.plot(list(range(1, N_iter_cavi1 + 1)), elbo1_cavi)
plt.title("ELBO plot")
plt.xlabel("Iterations")
plt.ylabel("ELBO")
plt.show()
```



### CASE 2

#### Small dataset

```
In [ ]: np.random.seed(0)

# Data simulation parameters
D2 = 1000
K2 = 50
K2 = 3
W2 = 10
eta_sim2 = np.ones(W2)
alpha_sim2 = np.ones(K2)

w2, z2, theta2, beta2 = generate_data(D2, N2, K2, W2, eta_sim2, alpha_sim2)

# Inference parameters
N_iter_cavi2 = 100
N_iter_svi2 = 100
eta_prior2 = np.ones(W2) * 1.
alpha_prior2 = np.ones(K2) * 1.
S2 = 100 # batch size

start_cavi2 = time.time()
phi_out2_cavi, gamma_out2_cavi, lambda_out2_cavi, elbo2_cavi = CAVI_algorithm(w2, K2, N_iter_cavi2, eta_prior2, alpha_prior2)
end_cavi2 = time.time()

start_svi2 = time.time()
phi_out2_svi, gamma_out2_svi, lambda_out2_svi, elbo2_svi = SVI_algorithm(w2, K2, S2, N_iter_svi2, eta_prior2, alpha_prior2)
end_svi2 = time.time()

final_phi2_cavi = phi_out2_cavi[-1]
final_gamma2_cavi = gamma_out2_cavi[-1]
final_lambda2_cavi = lambda_out2_cavi[-1]
final_phi2_svi = phi_out2_svi[-1]
final_gamma2_svi = gamma_out2_svi[-1]
final_lambda2_svi = lambda_out2_svi[-1]
```

#### Evaluation

Do not expect perfect results in terms of expectations being identical to the "true" theta and beta. Do not expect the ELBO plot of your SVI alg to be the same as the CAVI alg. However, it should increase and be in the same ball park as that of the CAVI alg.

```
In [ ]: np.set_printoptions(formatter={"float": lambda x: "{0:0.3f}".format(x)})
print(f"----- Recall label switching - compare E[theta] and true theta and check for label switching -----")
print(f"Final E[theta] of doc 0 SVI: (final_gamma2_svi[0] / np.sum(final_gamma2_svi[0], axis=0, keepdims=True)")
print(f"True theta of doc 0: (theta0[0], precision)")

print(f"----- Recall label switching - e.g. E[beta_0] could be fit to true theta.1. -----")
print(f"Final E[beta] k=0: (final_lambda2_svi[0, :]) / np.sum(final_lambda2_svi[0, :]), axis=-1, keepdims=True)")
print(f"Final E[beta] k=1: (final_lambda2_svi[1, :]) / np.sum(final_lambda2_svi[1, :]), axis=-1, keepdims=True)")
print(f"True beta k=0: (beta0[0, :]) / np.sum(beta0[0, :]), axis=-1, keepdims=True)")
print(f"True beta k=1: (beta0[1, :]) / np.sum(beta0[1, :]), axis=-1, keepdims=True)")

----- Recall label switching - compare E[theta] and true theta and check for label switching -----
E[theta] of doc 0 SVI: [0.410 0.590 0.090]
E[theta] of doc 0 CAVI: [0.238 0.338 0.424]
True theta of doc 0: [0.128 0.610 0.053]
----- Recall label switching - e.g. E[beta_0] could be fit to true theta.1. -----
E[beta] k=0: [0.142 0.858 0.083 0.145 0.027 0.058 0.018 0.183 0.161 0.093]
E[beta] k=1: [0.152 0.858 0.087 0.144 0.023 0.067 0.021 0.181 0.163 0.096]
E[beta] k=2: [0.164 0.866 0.083 0.144 0.030 0.069 0.021 0.184 0.177 0.071]
True beta k=0: [0.067 0.185 0.077 0.060 0.046 0.087 0.040 0.186 0.177 0.040]
True beta k=1: [0.138 0.867 0.074 0.138 0.087 0.068 0.002 0.158 0.134 0.181]
True beta k=2: [0.295 0.123 0.047 0.116 0.010 0.078 0.012 0.222 0.057 0.041]
Time CAVI: 0.015250509943
Time SVI: 44.14759862121582
```

```
In [ ]: plt.plot(list(range(1, N_iter_cavi2 + 1)), elbo2_svi(np.arange(0, N_iter_svi2, int(N_iter_svi2 / N_iter_cavi2))))
plt.plot(list(range(1, N_iter_cavi2 + 1)), elbo2_cavi)
plt.title("ELBO plot")
plt.xlabel("Iterations")
plt.ylabel("ELBO")
plt.show()
```



### CASE 3

#### Medium small dataset, one iteration for analysis.

```
In [ ]: np.random.seed(0)

# Data simulation parameters
D3 = 10**4
K3 = 500
K3 = 5
W3 = 10
eta_sim3 = np.ones(W3)
alpha_sim3 = np.ones(K3)

w3, z3, theta3, beta3 = generate_data_torch(D3, N3, K3, W3, eta_sim3, alpha_sim3)

# Inference parameters
N_iter3 = 1
eta_prior3 = np.ones(W3) * 1.
alpha_prior3 = np.ones(K3) * 1.
S3 = 100 # batch size

start_cavi3 = time.time()
phi_out3_cavi, gamma_out3_cavi, lambda_out3_cavi, elbo3_cavi = CAVI_algorithm(w3, K3, N_iter3, eta_prior3, alpha_prior3)
end_cavi3 = time.time()

start_svi3 = time.time()
phi_out3_svi, gamma_out3_svi, lambda_out3_svi, elbo3_svi = SVI_algorithm(w3, K3, S3, N_iter3, eta_prior3, alpha_prior3)
end_svi3 = time.time()

final_phi3_cavi = phi_out3_cavi[-1]
final_gamma3_cavi = gamma_out3_cavi[-1]
final_lambda3_cavi = lambda_out3_cavi[-1]
final_phi3_svi = phi_out3_svi[-1]
final_gamma3_svi = gamma_out3_svi[-1]
final_lambda3_svi = lambda_out3_svi[-1]
```

```
In [ ]: print(f"Examine per iteration run time.")
print(f"Time SVI: (end_svi3 - start_svi3)")
print(f"Time CAVI: (end_cavi3 - start_cavi3)")
Examine per iteration run time:
Time SVI: 2.00354525188374
```



